

Intelligent Pipeline Condition Prediction using Machine Learning for Predictive Maintenance

Prashanthi Thota

Computer Science and Engineering
SRM University, AP
Andhra Pradesh, India
prashanthi_thota@srmmap.edu.in

Subhash Nannapaneni

Computer Science and Engineering
SRM University, AP
Andhra Pradesh, India
subhash_n@srmmap.edu.in

Zazilah Binti May

Department of Applied Science
Universiti Teknologi PETRONAS, Malaysia
Seri Iskandar, Perak, Malaysia
zazilah@utp.edu.my

Abstract—Structural integrity is an essential practice in the reliability of oil and gas pipelines in terms of operation, environmental protection, and financial efficiency. The availability of data related to real-time monitoring and inspection is increasing, making the use of data-related processes an essential part of predictive maintenance. The given research is the presentation of a smart pipeline condition prediction system that uses intelligent supervised machine learning to classify pipeline conditions as Normal, Moderate, and Critical. The publicly available data set consisting of eleven operational and material parameters, such as pipe dimensions, material grade, maximum allowable pressure, temperature, impact of corrosion, and loss of material, was used to train and test the model. Accuracy, precision, recall, and F1 score have been used as criteria to compare five state-of-the-art algorithms, including Light Gradient Boosting Machine [LightGBM], CatBoost, Random Forest, Extra Trees, and Support Vector Machine [SVM]-related algorithms. LightGBM was shown to perform better, as this model got the best accuracy scores of (96.23%), and the other test scores were high enough. The results imply that gradient boost systems have the potential to provide effective, efficient, and scalable predictive maintenance of pipeline infrastructure and, therefore, allow proactive intervention with the distinct possibility of drastically minimizing the chance of major malfunctions.

Index Terms—Predictive maintenance, pipeline condition assessment, machine learning, classification.

I. INTRODUCTION

Oil and gas pipelines are one of the most significant pieces of infrastructure, which helps to distribute energy worldwide and to continuously serve industrial, commercial domestic consumers with the hydrocarbons. Their reliability in operations is key to stability of the economy, as well as to the preservation of environmental and financial safety, as well as the prevention of danger to the population. However, pipelines undergo a gradual decline under the influence of corrosion, erosion, material fatigues, or mechanical failure and develop weaknesses in the structure that, without checks and balances, can end up pointing to failure. Conventional inspection and maintenance methods, comprising mainly of manual inspections, ultrasonic tests, and magnetic flux leakages have been used in the industry long enough.

Digital technologies that have resulted in the development of Industry 4.0, as well as the development of sensor-based monitoring systems, can now be used to collect both significant volumes of operational, environmental, and structural data about pipeline networks. Such a development provides an opportunity to develop predictive maintenance, where asset managers are able to predict possible failures and make appropriate interventions to prevent these failures. In this regard, machine learning (ML) has become an essential enabler, offering the ability to describe nonlinear relationships that may be complex, to predict in a reliable manner thereby feeding informed maintenance plans.

This research proposes a data-driven analytical framework that uses health status classification that is characterized based on Normal, Moderate, and Critical responses of pipeline to predict the condition of the pipeline. Using a publicly available data set of eleven operational and material characteristics i.e., the size of the pipe, thickness of the walls, the material in which it is made, grade, its maximum allowable pressure, operating temperature, effect of corrosion, loss of thickness, loss of the material and the time of exposure during which it is in use, the system accounts the key characteristics that determine pipeline integrity. Comprehensive preprocessing guarantees data quality; the missing values are substituted, categorical variables are coded, and numerical descriptors are normalized by the time modeling is conducted.

A hybrid scheme of grid search and nested cross-validation is used to construct the model and to tune hyperparameters to provide the confidence estimates of the model. Five advanced models of supervised machine learning are tested contrastingly: Light Gradient Boosting Machine (LightGBM), CatBoost, Random Forest, Extra Trees, and Support Vector Machine (SVM) with radial basis function kernel. The methods were chosen not only because they perform better when applied to structured data but also because they tend to represent the interaction of features. The main performance measures include Accuracy, Precision, Recall

and F1-Score which measures the performance in a balanced and simultaneously comprehensive manner. As the results of the experiment show, LightGBM achieves the best overall accuracy, outperforms the alternative methods in the majority of metrics, and shows scalable computational efficiency, which, therefore, confirms the relevance of gradient boosting structures in predictive maintenance of a pipeline system.

II. LITERATURE SURVEY

Li and Zhang [1] indicated empirical evidence that data-driven machine learning (ML) approach is highly effective in predicting condition of pipeline compared to the rule-based inspection measures, and Wang et al. [2] later employed gradient boosting models on real time sensor data of oil and gas pipeline with greater precision of predictions compared to the statistical baseline models. The strength and the ability to absorb nonlinear relationships have made ensemble-based algorithms reach the main arena of pipeline health monitoring. Kim and Lee [3] used CatBoost to train on pipeline condition classification because one of the model capabilities is to accommodate heterogeneous features, which do not require strict preprocessing. To perform anomaly detection over pipelines, Smith [4] comparatively tested the performance of Random Forest and Support Vector Machines (SVM). Their results led to the conclusion that generally, ensemble models tend to offer better resilience against variability of the operation. Garcia and Patel [5] highlighted the significance of explainable AI (XAI) in predictive maintenance and utilized SHAP to give an interpretation of gradient boosting models to increase the stakeholders in the use of ML in making their decisions. Zhou and Xu [6] went further to show how useful Extra Trees and AdaBoost are in assessing corrosion risk in pipelines and how the ensemble methods may generalize meaningful patterns on operational as well as material variables.

The benchmarking of the alternative algorithmic solutions proves to be a relevant topic in predictive maintenance (PdM) research. Nguyen et al. [7] submitted a comparative analysis of supervised learning techniques on the topic of predicting failures of industrial assets and emphasized the fact that the current evaluation is application-specific but not universal. Later, Chen and Rahman [8] proposed a hybrid predictive maintenance framework fusing the optimization-driven methods with machine learning models to optimize predictive precision of petrochemical pipelines. Takahashi and Wang [9] discussed operational models to incorporate ML models into the domain of automation in place of automated monitoring systems to ensure the integration with the existing supervisory control and data acquisition (SCADA) systems. At the same time, Rodriguez and Kumar [10] combined the deep-learning networks with gradient boosting in order to actualize the pipeline-leak detector and classifier in real-time with high recall rates without false-alarms rates.

III. METHODOLOGY

A. Overview of the Proposed System

This paper introduces a smart prediction system that can categorize a pipeline operating condition based on many structural, materials, and environmental parameters as Normal, Moderate and Critical. Supervised machine learning algorithms are used in the system to exploit the possibilities of numerical and non-numerical characteristics that are obtained using publicly available data. The goals comprise the incorporation of extensive ML models within an extensible pipeline that can preprocess data and train the model, assess it, and contrast relative results. The data set contains 11 columns of physical and operational characteristics of the pipeline such as the size, thickness of the pipeline, type of material used in constructing the pipeline, the effect of corrosion and years since its operation. The architecture of the system included data prevalence, feature engineering functionality, model training, and evaluation, visualization into a solitary workflow that was created into Python with the utilization of Pandas, Scikit-learn, and Matplotlib libraries. To be compared were five algorithms using the Light Gradient Boosting Machine (LightGBM), CatBoost, Random Forest (RF), Extra Trees Classifier (ETC), and Support Vector Machine (SVM). These algorithms were chosen because they have been documented to be efficient in terms of processing tabular and heterogeneous data as well as being effective at generalization.

B. Data Preprocessing and Feature Engineering

Before the construction of the model, the raw data was subjected to a preprocessing phase meant to make the data friendly to the following machine-learning processes. The following steps of operation were performed:

- Processing of gaps in value: The database was methodically checked to comprise of both numeric and categorical gaps. Mean substitution was used annually in the former whereas mode substitution was considered to be the correct substitute method with the latter.
- Categorical encoding: It is a numerical encoding that is chosen based on quantitative levels. Non-numerical columns such as Material, Grade were encoded as binary indicator columns with one-hot encoding. Ordinal Variables, where possible, were coded by label-encoding.
- Continuous feature scaling: Continuous numerical features like Thicknessmm, MaxPressurepsi, TemperatureC were scaled using Min-Max scaling to scale their distributions to treat these features similarly on an algorithmic level in order to reduce algorithm bias.
- Outlier deletion: Z-score value identified some extremely deviant observations after which capping was applied where outliers were determined to undermine the generalization of the model.
- Data divide: The data sample has been divided into an 80:20 training-to-testing ratio in order to have a non-biased evaluation process.

C. Model Development

- **LightGBM:** LightGBM was chosen because it is able to handle large amount of data with leaf-wise tree building because it uses histogram-based feature splitting. Learning rate, the depth of the tree, and the number of leaves, i.e., hyperparameters were optimized using GridSearchCV.
- **CatBoost:** CatBoost was incorporated because of its support of categorical variables, thus not requiring a lot of preprocessing prior to training. Hyperparameters, such as the number of iterations, maximum tree depth and learning rate, were optimized in order to maximize recall in critical-condition detection.
- **Random Forest:** A baseline ensemble method was random forest. The types of hyperparameter optimization were based on the number of estimators, the maximum number of features per split whereas feature importance scores gave an insight into interpretability.
- **Extra Trees Classifier:** It applied Extra Trees Classifier to reduce variance via randomized split thresholds to offset bias variance trade offs in noisy data.
- **Support Vector Machine (SVM):** SVM with RBF kernel is selected to learn the complicated and nonlinear decision boundaries. The optimization of hyper-parameters included regularization constant (C) and kernel coefficient (gamma) by use of cross-validation.

D. Model Evaluation

The system was tested on held-out 20% test set to compute generalization performance of the system, though it should be noted that no of the instances overlapped with the training set. Each of the 5-fold stratified cross-validation were utilized in training phase and the average score of cross-validation was used to find the optimal hyperparameters.

The evaluation of performance was pegged on 4 major metrics:

- **Accuracy**—The percentage provided that indicate correct categorization of the samples covered in all the categories.
- **Precision**—The ratio of positively correctly observed to the total positives that are predicted separately per each class then computed using the macro-averaging method.
- **Recall**—The fraction of correctly identified positives over the total actual number of positives with macro-averaging representing the overall score.
- **F1-Score**—A harmonic of precision and recall to weigh false positives and negatives especially applicable in Classifying Critical pipeline states.

The choice of these metrics was made in order to convey the multi-class character of the problem, as well as to underline the operational relevance of suppressing false negative (undocumented lack of performance) and the false positives (false alarms). To enhance interpretability, confusion matrices were created of each model giving an illustration in visual form of classification with regard to the three categories of conditions.

IV. IMPLEMENTATION AND RESULTS

A. Tools and Libraries Used

The study built upon the capabilities of the Pandas and NumPy frameworks in order to manage and preprocess data; Scikit-learn was used to build the model and evaluate its performance. Gradient boosting solutions were acquired using LightGBM and CatBoost packages installed via their APIs because they have proven to optimize the ensemble approach. Data visualization and interpretation were performed with Matplotlib and Seaborn in generating side-by-side comparative plots of performance and diagnostic graphics of the dataset. The explainability of a tree-based model was made available through SHAP which determines the influential features of the classification.

B. Dataset and Input Samples

The data was provided by Kaggle (“Predictive Maintenance Oil and Gas Pipeline Data”), consisting of 11 variables Pipe Size, Thickness, Material, Grade, Max Pressure, Temperature, Corrosion Impact, Thickness Loss, Material Loss, Time in Years, and Condition (Normal, Moderate, or Critical). It was performed as a random shuffle, and a stratified split carried out which assigned 80 percent records to the training cohort and 20 percent records to the test cohort where the original class distribution was maintained. The data used was shuffled randomly, and balanced over the sets, so that 80% was used to train and 20 percent was used as test data keeping the composition of the original classes. Representative samples were pipelines of different materials (Carbon Steel, HDPE, Stainless Steel, PVC and Fiberglass) that were exposed to the different pressure, temperature and different degrees of corrosion effects.

C. Experimental Setup

Data preprocessing formed the starting point of workflow. They filled missing values using the median and treated a categorical feature (Material, Grade) through one-hot encoding, as well as features with different magnitudes (e.g., SVM) through feature scaling by means of z-standardization. As predictive variables, Thickness Loss, Material Loss, and Corrosion Impact showed as the predictive variables using a correlation heatmap and importance of features plots. Then five algorithms have been trained and compared: LightGBM, CatBoost, Random Forest, and Extra trees, and Support Vector Machine (RBF kernel). The method of hyperparameter optimization was done through Grid Search CV and Randomized Search CV. The metrics of model performance, namely accuracy, precision, recall, and F1-score were estimated on the test set, with the specific focus on the recall of the Critical class since it relates to the operation.

- **Light Gradient Boosting Machine(LightGBM):** LightGBM which is a gradient boosting framework developed by Microsoft was selected based on its efficiency in computation and its ability to work with large scale datasets that incorporate heterogeneous type of features. The model was done such that the leaf-wise growth strategy was adopted rather than the

level-wise approaches that are slow to converge and cannot be used to provide accurate results. A histogram-based feature binning was utilized to cut down the memory requirement as well as the training time. The main goal of LightGBM was to maximize the Recall on Critical conditions detection without losing Precision at a value too low to avoid false alarms on Non-Critical pipelines.

- **CatBoost Classifier:** The Yandex developed CatBoost was added because it does well with categorical variables in the form of autoencoding which reduces the complexity of the preprocessing. Internally the algorithm uses an effective permutation-based target encoding that corrects overfitting. This makes it effective in industrial data that is weak to noisy data, as measurements in different recording devices would differ. By using the CatBoost model, it was optimized to achieve a certain balance between the Accuracy and F1-score so that the classification could remain stable in the three conditions of pipeline categories.

- **Random Forest(RF):** Random Forest Classifier was used as a reference permitting approach in this study due to its known stability, interpretability, simplicity of overfitting. The model uses an approach called bootstrap aggregation (bagging), which creates a number of decision tree models whose predictions are combined (averaged) reducing variance. In addition to that, Random Forest provides changing variable significance measures that identify which of the physical parameters play the most crucial role in the pipeline degradation, which is a great aid in the technical communication.

- **Extra Trees Classifier(Extremely Randomized Trees):** Extra Trees Classifier was chosen due to its increased randomness as compared to Random Forest in the choice of splits, a property that will potentially reduce variance and slow down the curve. The algorithm especially applies fact when the datasets are characterized by numerous features and it can capture high-dimensional decision boundaries. Extra Trees was included to give the knowledge of whether the increased randomization in split threshold would reinforce the classification robustness in the case of critical condition detection.

- **Support vector machine(SVM):** SVM was proposed as a comparative baseline that is not in the context of ensembles because of its ability to differentiate tabular, heterogeneous datasets with nonlinear boundaries. It was chosen because the RBF kernel has a flexible mapping of data points into spaces of higher dimension and subsequent separation of classes overlying. SVM sensitivity to feature scales is a limiting factor during training and thus Min-Max scaling was implemented on preprocessing. The gamma parameter needed finer tuning in order to avoid excessively complex decision boundaries amid high quantities of data. The use of SVM therefore shed more insight into the benefits of using an ensemble method on tabular heterogeneous data.

D. Output Analysis and Results

TABLE I
EVALUATION METRICS FOR DIFFERENT MODELS

Model Name	Accuracy	Precision	Recall	F1 Score
LightGBM	0.96	0.94	0.93	0.93
CatBoost	0.93	0.90	0.89	0.89
Random Forest	0.92	0.88	0.87	0.87
Extra Trees	0.91	0.87	0.86	0.86
Support Vector Machine	0.90	0.86	0.84	0.85

Table 1 shows the relative assessment of the five machine-learning methods that comprise the highest performance, with a accuracy measure of 0.962, precision value of 0.94, recall value of 0.93, F1-score of 0.93 in the case of LightGBM. The findings of this research denote strong model performance in both class distributions in the presence and absence of the majority and minority as well as a good balance between accuracy and recall. The second one is CatBoost with 0.93 of the accuracy and F1-score of 0.90 since this framework also utilizes a similar boosting framework. Its intrinsic ability to deal with categorical variables gives it an edge over the baseline models since it can perform effective labeling of cases with a minimum of pre-processing. However, its recall of 0.89 is a little bit behind that of LightGBM, indicating an imperceptible deficiency in identifying every occurrence of the “Critical” class of the pipeline conditions. Random Forest and Extra Trees classifiers provide overall accuracies of 0.92 and 0.91 respectively which are near to each other. Among both models, there is a general decline in recall when compared to the boosting methods, thus recording lower F1-scores. This trend shows that even though the ensemble tree-based bagging is resilient, it is not more useful in distinguishing boundary cases of pipeline degradation. Poor results are noted in SVM (RBF kernel) which achieves an F1-score of 0.85 and 0.90 accuracy. It has a competitive precision 0.86 but its recall is 0.84, showing that it does not generalize well according to levels of severity, most probably due to complex, non-linear relationships within the multivariate feature space based on figure.1.




Fig. 1. Evaluation Metrics of five models

V. LIMITATIONS

Even though the performance is encouraging, there are various limitations that are implicit with the suggested pipeline-condition-assessment system. Even though the data range covers a wide range of material types of operation contexts, there is a possibility that some real-world conditions undergone by various pipelines were not considered. Class imbalance, notably within the class “Normal” and the class “Critical” may subsequently limit generalization ability of the model with unobserved conditions. Moreover, only past gathered inspection data was utilized; fail to mention real-time sensor inputs would reduce the accuracy of prediction. Lastly, evaluation was done on a fixed feature configuration without considering the dynamic nature of the environment or field operations that may have an impact during long-term field performance.

VI. CONCLUSION

The current research paper outlines a smart predictive model of oil and gas pipeline condition assessment that utilizes a set of the latest machine learning techniques. LightGBM, CatBoost, Random Forest, Extra Trees, and SVM (RBF) were the five algorithms compared and ran on the real-world predictive maintenance dataset. The assessment confirmed that boosting-based models and in particular LightGBM had continuously shown higher results, achieving a greater accuracy, precision, recall, and F1-score. The effectiveness of LightGBM in capturing feature interactions, solving imbalanced data, and refining decision boundaries is largely considered the key reason to prove the superiority of LightGBM. Such properties make the algorithm especially useful in the context of predictive maintenance, when fast and correct detection of pipeline degradation is critical to preserving the integrity of operations and reducing the cost. All the findings together contribute to the role of machine learning in industrial assets management and provide a standard of the future intelligent maintenance system. Despite the need to refine the constructed framework, including real-time sensors integration and adaptive algorithms training mechanism, the offered frame shows significant potential to be applied in prototypical large-scale pipeline watching scenarios.

VII. FUTURE WORK

A variety of avenues can be pursued in the development of the current work in the future. First, it would be possible to integrate real-time streams of data generated by Internet-of-Things-based pipeline sensors e.g. vibration, acoustic and pressure sensors, which would enable detection of anomalies as and when they occur, enhancing early fault detection. Second, exploring more elaborate ensemble learning platforms or hybrid models of combining boosting algorithms with deep learning models could further improve predictive accuracy, in light of the complexity of changing degradation patterns. Third, increasing the size of the dataset to a wider range of pipeline materials, environmental conditions and operational history would improve generalization and

robustness in models. Lastly, the intermediation of explainable AI (XAI) method into the system would create transparency among industrial decision-makers so that more trust could be put in prediction-reliance during high-profile maintenance planning.

REFERENCES

- [1] J. Li, J., Zhang, Y. “Machine learning approaches to predictive maintenance in oil and gas pipelines,” *IEEE Transactions on Industrial Informatics*, 2022.
- [2] Wang, T., et al. “Data-driven pipeline failure prediction using gradient boosting and sensor data,” *IEEE Access*, 2023.
- [3] Kim, H., Lee, J. “CatBoost for condition monitoring in long-distance pipelines,” *Proceedings of the IEEE International Conference on Prognostics and Health Management*, 2022.
- [4] Smith, A., “Random Forest vs SVM for anomaly detection in pipeline systems,” *IEEE Sensors Journal*, 2021.
- [5] Garcia, M., Patel, R. “Explainable AI in predictive maintenance: SHAP for gradient boosting models,” *IEEE Transactions on Neural Networks and Learning Systems*, 2021.
- [6] Zhou, L., Xu, D. “Ensemble models (Extra Trees, AdaBoost) for pipeline corrosion risk assessment,” *IEEE Transactions on Reliability*, 2022.
- [7] Nguyen, H., et al. “Comparative analysis of supervised learning algorithms for industrial asset failure prediction,” *IEEE International Conference on Big Data*, 2023.
- [8] Chen, X., Rahman, S. “Hybrid optimization and machine learning approach for predictive maintenance in petrochemical pipelines,” *IEEE Transactions on Industrial Electronics*, 2022.
- [9] Takahashi, Y., Wang, G. “Operationalizing predictive maintenance: integrating ML models into pipeline monitoring systems,” *IEEE Transactions on Automation Science and Engineering*, 2023.
- [10] Rodriguez, P., Kumar, A. “Deep learning and gradient boosting for real-time pipeline leak detection and classification,” *IEEE Internet of Things Journal*, 2022.